

# Distributed Systems

## Exercise 1: Sockets / Java NIO



# Agenda

Organizational matters

Follow-up: idea behind sockets

Java New I/O

- motivation
- buffer and channels
- transformation characters  $\Leftrightarrow$  bytes
- asynchronous I/O using the selector approach

Notes regarding exercise sheet 1

- SMTP overview
- notes and tips



## Organizational Matters

Tutors:

- Gabriel Behrendt: [gabriel.behrendt@campus.tu-berlin.de](mailto:gabriel.behrendt@campus.tu-berlin.de)
- Aurel Isaak Weinhold: [a.weinhold@campus.tu-berlin.de](mailto:a.weinhold@campus.tu-berlin.de)

Material published on the ISIS website:

- lecture slides
- guidelines & tools (e.g. code templates)

Exam authorization:

- There will be 4 exercise sheets which will be evaluated

Exercise schedule published on the ISIS website



# Organizational matters

## Procedure

- 4 assignments throughout the semester
- Work in teams of 5 students
- One exercise sheet per assignment
- Introductory material at beginning of a assignment
- Tutorial 2 weeks later for resolving problems and asking questions



# Appointments

Date	
24.04.22	Group Selection
25.04.22	Publish exercise 1
06.05.22	Tutorial exercise 1
13.05.22	Hand-in exercise 1
16.05.22	Publish exercise 2
27.05.22	Tutorial exercise 2
03.06.22	Hand-in exercise 2
06.06.22	Publish exercise 3
14.06.22 (!!Dienstag!!)	Tutorial exercise 3
24.06.22	Hand-in exercise 3
27.06.22	Publish exercise 4
08.07.22	Tutorial exercise 4
15.07.22	Hand-in exercise 4



## Idea behind Sockets

Definition of distributed systems given during the lecture:

A Distributed System is one in which hardware or software components are located at networked computers and communicate and coordinate their actions only by passing messages.

© Colouris

Problem: heterogeneous hard- and software components

- How to exchange messages between them?
- The operating system has to provide a uniform interface to the applications


## Solution: Sockets

Uniform interface used for network- and inter-process communication (IPC)

- Originally developed for Berkeley Unix (1981)
- Meanwhile supported among most of available platforms

Represent end-point of a communication channel

Each endpoint of the connection corresponds to one process






# What is a socket?

Process requests a socket from the OS using:

- address and port used to access the process inside the network
- protocol to transmit data

Combination of address/port/protocol uniquely identifies a process in the network







# Sockets and Protocols

Sockets can operate using different transmission protocols

Protocol selection has impact on the characteristics of the communication

Typical classification of transmission protocols:

- connectionless communication (e.g. UDP)
- connection oriented communication (e.g. TCP)
- (local communication)
- (proprietary communication protocols)



# Basic Socket Operations

Basic operations when working with sockets

- Applicable for both connectionless and connection oriented protocols
  - `socket()`: creates a socket
  - `bind()`: bind process to address/port (server only)
  - `send()`: send data
  - `recv()`: receive data
  - `close()`: close socket
- Applicable for connection oriented protocols only
  - `connect()`: establish connection (client only)
  - `listen()`: prepare socket to accept new connections (server only)
  - `accept()`: wait for new connections (server only)



## What is Java New I/O?

New API for Input-/Output (I/O) operations

- e.g. file-I/O, network-I/O

Benefits for I/O intensive applications (usually servers)

- faster execution of I/O operations
- higher scalability in case of many concurrent clients

Introduced 2002 along with Java 1.4

- updated with Java SE7





## Why Java New I/O?

Many server applications can be characterized by:

- high requirements regarding I/O performance
  - e.g. File Server, Web Server, Mail Server, ...
- many concurrent client requests
  - several data streams need to be maintained in parallel
- heterogeneous clients
  - possibly different data formats
  - possibly different standards for character encoding
  - frequent source of error

## Why Java New I/O?

former Java-abstraction for I/O is based on streams

- Advantages: intuitive, combinable

```
Socket client = server.accept();
InputStream inputStream = client.getInputStream();
while((int i = inputStream.read()) != -1) {
    /*... do something... */
}
```

- Disadvantages:
  - Byte-per-Byte-processing slow/inefficient (performance)
  - read-/write-operations blocking → many threads required (performance, scalability)

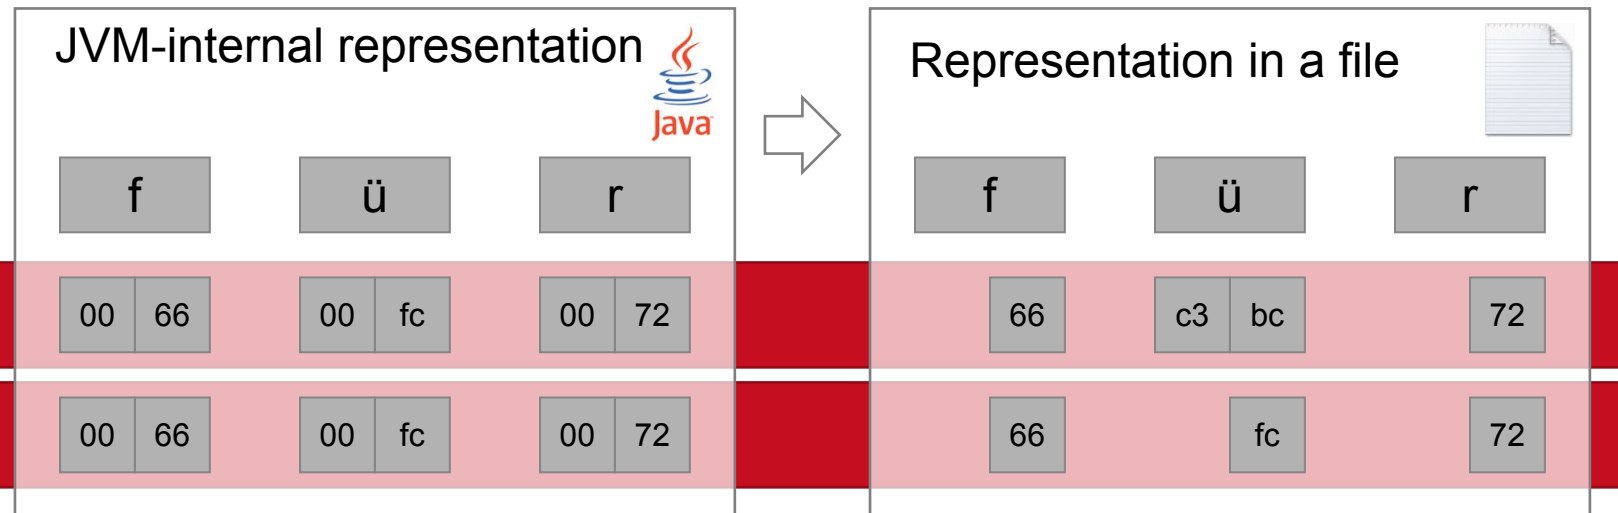
# Why Java New I/O?

Stream oriented I/O classes distinguish between

- Byte streams: Byte (8 Bit)
- Character streams: UTF-16 encoded double bytes (16 Bit)

Encoding Character stream  $\Leftrightarrow$  Byte stream frequent source of errors

- problem: Java uses default-charset of operating system



# Buffers and Channels


Java NIO uses buffers instead of Byte-per-Byte

- Idea: memory backing buffer is allocated in regions the operating system uses for its I/O operations
- => no CPU overhead required due to copy operations

Channels are bidirectional interfaces for I/O

- Incoming data are written into the buffer
- Outgoing data are read from the buffer

View of the developer







# Introduction to Buffers

Buffer classes of NIO have three important properties

- position: position for next read/write operation
- limit: position until which read/write is allowed
- capacity: capacity of the buffer (static)




ByteBuffer of length 6,  
initial state




ByteBuffer of length 6,  
state after inserting two bytes

## Important buffer operations

- flip-operation: limit = position, then position = 0



- clear-Operation: position = 0, then limit = capacity



```
byte [] test = {'T','e','s','t'};
```

```
ByteBuffer buf = ByteBuffer.allocate(8);
```

```
buf.put(test);
```

```
buf.flip();
```

```
FileOutputStream f;
```

```
f = new FileOutputStream("test.txt");
```

```
FileChannel ch = f.getChannel();
```

```
ch.write(buf);
```

```
ch.close();
```

```
buf.clear();
```

position: 0  
limit: 8  
capacity: 8

position: 4  
limit: 8  
capacity: 8

position: 0  
limit: 4  
capacity: 8

position: 4  
limit: 4  
capacity: 8

position: 0  
limit: 8  
capacity: 8

## NIO buffer class hierarchy

Im Paket `java.nio`

Abstract base class

**Buffer**

Sub classes for concrete data types (inherit from Buffer)

**ByteBuffer**

**CharBuffer**

**FloatBuffer**

**IntBuffer**

**LongBuffer**

**ShortBuffer**



## Methods for Channels

Important classes that implement the channel interface

- `FileChannel`
- `SocketChannel`
- `ServerSocketChannel`

Important methods used with channels

- `write(ByteBuffer src)`: writes `src` into channel
- `read(ByteBuffer dst)`: writes content of channel into `dst`
- `close()`: closes channel



## Transformation

### Characters → Bytes

Java NIO introduces explicit classes to encode/decode between strings and byte arrays

Developer can define charset

```
String s = "Verteilte Systeme";  
Charset messageCharset = null;  
  
try {  
    messageCharset = Charset.forName("US-ASCII");  
} catch (UnsupportedCharsetException uce) {...}  
  
byte [] b = s.getBytes(messageCharset);  
byteBuffer.put(b);
```

Transformation

Bytes → Character

```
CharsetDecoder decoder = messageCharset.newDecoder();  
  
try {  
    CharBuffer charBuf = decoder.decode(byteBuffer);  
} catch (CharacterCodingException e) {...}  
  
String s = charBuf.toString();
```

`CharBuffer` implements `CharSequence` interface

- Explicit conversion into `String` often unnecessary
- e.g. matching of regular expressions





# Asynchronous I/O Using the Selector Approach

I/O-API before Java NIO was solely synchronous

- blocking operations
- one thread per connection
- results in loss of performance and bad scalability

Java NIO offers asynchronous I/O API if operating system support available

- application triggers I/O operation and keeps on executing
- one thread is able to maintain several connections



# The Selector Class

**Selector**-class serves as a registry for channels

- Channel has a set of supported events
- Developer defines which events are of interest

Possible events:

- OP\_CONNECT: connection established successfully (client)
- OP\_ACCEPT: new connection established by client (server)
- OP\_READ: Data are ready to read from the channel
- OP\_WRITE: Channel is ready to write data

## Example: Create/Register ServerSocket

```
Selector selector = Selector.open();

ServerSocketChannel servSock = ServerSocketChannel.open();
servSock.configureBlocking(false);
servSock.socket().bind(new InetSocketAddress(6332));

servSock.register(selector, SelectionKey.OP_ACCEPT);
```

`ServerSocketChannel` encapsulates `ServerSocket`

- Socket listens on port 6332
- Notification about event `OP_ACCEPT` requested
- Note: no exception handling shown → required for real applications!


# Notifications About Occurred Events

## Interest-Set

- Events related to a channel we are interested in
- Define by the parameters of the `register()` method

## Ready-Set

- Registered events that have occurred
- Subset of the interest-set of a channel



Are there channels with non-empty ready-set?

- Query using `select()` method of the `Selector` class

## Example: Notification

```
...  
while(true) {  
  
    if(selector.select() == 0) /* blocking */  
        continue;  
  
    Set<SelectionKey> selectedKeys = selector.selectedKeys();  
    Iterator<SelectionKey> iter = selectedKeys.iterator();  
  
    while(iter.hasNext()) {  
  
        SelectionKey key = iter.next();  
        /* check ready set of channel */  
        iter.remove();  
    }  
}
```



## The Class `SelectionKey`

`SelectionKey` is the key to a channel with non-empty Ready-Set

- Important methods to check the state of a `SelectionKey`
  - `isReadable()`: Tests if data is ready for reading
  - `isWritable()`: Tests if channel is ready to write data
  - `isAcceptable()`: Tests if connection is ready to be accepted (server)
  - `isConnectable()`: Tests if connection is ready to be established (client)
  - `cancel()`: Unregister selector with channel
  - `channel()`: Access to the corresponding channel
  - `attach(Object ob)`: Attach additional data to the channel (e.g. a session object identifying the client)

## Example: Event Query

```
...
Iterator<SelectionKey> iter = selectedKeys.iterator();

while(iter.hasNext()) {

    SelectionKey key = iter.next();

    if(key.isAcceptable()) {
        /*... do something ...*/
    }
    if(key.isReadable()) {
        /*... do something ...*/
    }

    iter.remove();
}
```



## Example: New Connection Accepted

```
...  
if(key.isAcceptable()) {  
  
    ServerSocketChannel sock =  
        (ServerSocketChannel) key.channel();  
  
    SocketChannel client = sock.accept();  
    client.configureBlocking(false);  
    client.register(selector, SelectionKey.OP_READ | SelectionKey.OP_WRITE);  
}  
...
```

New `SocketChannel` for incoming connection

- configured as non-blocking
- afterwards registered with the `Selector`

## Example: Read Data From the Channel

```
...  
if(key.isReadable()) {  
  
    ByteBuffer buf = ByteBuffer.allocate(1024);  
  
    SocketChannel channel = key.channel();  
  
    channel.read(buf);  
  
    buf.flip();  
  
    /*Further processing of data*/  
}  
...
```

## Overview SMTP

Simple Mail Transfer Protocol used to exchange e-mails

- Simple ASCII protocol based on the Request/Reply-pattern

Well defined protocol-handshake

- HELO <Hostname>
- MAIL FROM: <e-mail address of the sender>
- RCPT TO: <e-mail address of the recipient>
- DATA
- <content of the e-mail, terminated with \r\n.\r\n>
- QUIT
- (HELP)





## Notes Regarding Exercise 1

In order to test your server implementation, a test client is published on the ISIS website

- telnet-client helpful for testing first functionalities

Test client judges success of a request by inspection of the returned response code

The test client will randomly interrupt the protocol handshake with HELP commands!



## Tips & Additional Material

### Java API-Docs

- <https://docs.oracle.com/javase/8/docs/api> (first introduction of Java Streams)
- <https://docs.oracle.com/javase/10/docs/api/>

### Source code of the SMTP-client

### ISIS forums